

8th Vienna Deep Learning Meetup

Jan 17, 2017 @ TU Wien



Thomas Lidy



Jan Schlüter



Alex Schindler



8th Vienna Deep Learning Meetup

Agenda:

- Welcome
- Hot Topics and Latest News (Tom Lidy & Jan Schlüter)
- Coding Session with Keras:
Deep Learning Library for Python (Tom Lidy, Alex Schindler)
- Questions & Answers
- Discussion



Latest News

Hot Topics

a 5-10 min block at every meetup to briefly present
“trending topics”

Send us contributions (tom.lidy@gmail.com)
or come with slides to do a 5-10 min block yourself!



Apple starting to publish AI research?

“The company will begin sharing its research into artificial intelligence in a move that may help advance its efforts.”


<https://www.cnet.com/news/apple-is-going-start-publishing-its-research-into-ai/>



hardmaru
@hardmaru

Folgen

Apple will start publishing, according to @rsalakhu at #nips2016



RETWEETS
234

GEFÄLLT
401



06:07 - 6. Dez. 2016 aus Barcelona, Spanien



TensorFlow now supports Windows

before, it was only possible using Docker or VMs

- not a good way for TensorFlow to access a GPU



TensorFlow 0.12 adds native support for Windows:

<https://developers.googleblog.com/2016/11/tensorflow-0-12-adds-support-for-windows.html>

Beware: TensorFlow supports **only 64-bit Python 3.5** on Windows

Guide: How to install TensorFlow on Windows without Docker / VMs:

<http://www.netinstructions.com/how-to-install-tensorflow-on-windows-without-docker-or-virtual-machines/>

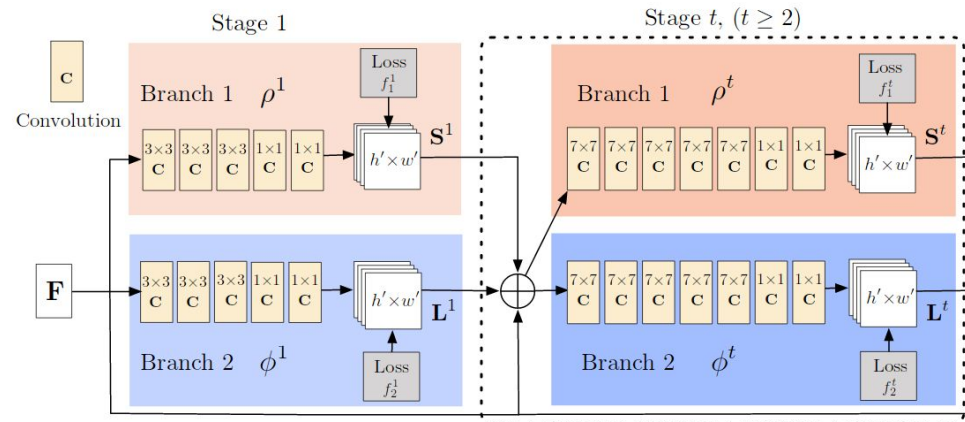
Thanks to by Dalibor Mitrovic for recommending this news contribution!



Realtime Multi-Person Pose Estimation



two-branch multi-stage CNN:



<https://www.youtube.com/watch?v=pW6nZXeWIGM&t=77s>
<https://arxiv.org/abs/1611.08050>
Code link under the video

*Recommended by
Alex Schindler*



Google open-sources DeepMind Lab



DeepMind Lab platform was used for AI to learn how to play computer games, navigate through mazes, etc.

One of the goals:

“Design intelligent agents capable of more-and-more sophisticated cognitive skills”

Now available to broader research community, as an opportunity to make further significant contributions.

Fully 3D game-like platform tailored for agent-based AI research, to learn: navigation, memory, 3D vision from a first person viewpoint, motor control, planning, strategy, time, etc.



<https://deepmind.com/blog/open-sourcing-deepmind-lab/>



OpenAI releases “Universe”

Universe: a software platform for measuring and training an AI's general intelligence for of games, websites and other applications.

allows an AI agent to **use a computer like a human**: looking at screen pixels and operating a virtual keyboard and mouse (through VNC)

“In April, we launched Gym, a toolkit for developing and comparing reinforcement learning algorithms. With Universe, any program can be turned into a Gym environment.”

“We must train AI systems on the full range of tasks we expect them to solve, and Universe lets us train a single agent on any task a human can complete with a computer.”



Universe includes 2600 Atari games, 1,000 Flash games, browser tasks and more
<https://openai.com/blog/universe/>



Uber launches AI lab

Uber acquired AI start-up Geometric Intelligence with 15 ML researchers and launched an AI lab in San Francisco

“as it seeks to use the technology to improve its driverless cars and delivery routes”

(first “self-driving” Uber cars now in SF, but with a “person monitoring the car”)

<https://www.ft.com/content/7ad79536-bb0c-11e6-8b45-b8b81dd5d080>

<https://newsroom.uber.com/ailabs/>



Practical Deep Learning For Coders

7 week course: <http://course.fast.ai>

- SGD, CNNs, RNNs
- Over-/underfitting
- Cats&Dogs, NLP, Embeddings
- Keras functional API
- ...

```
▶ Reducing overfitting [...]
▼ Resnet

In [4]: import resnet50; reload(resnet50)
        from resnet50 import Resnet50

In [5]: rn0 = Resnet50(include_top=False).model

In [7]: rn0.output_shape[1:]
(2048, 7, 7)


In [6]: batches = get_batches(path+'train', shuffle=False, batch_size=batch_size)
        val_batches = get_batches(path+'valid', batch_size=batch_size*2, shuffle=False)
        (val_classes, trn_classes, val_labels, trn_labels,
         val_filenames, filenames, test_filenames) = get_classes(path)
```

Thanks to by Dalibor Mitrovic for recommending this news contribution!



Gradient-based Hyperparameter Optimization through Reversible Learning

- Backpropagate through backpropagation to adjust hyperparameters (learning rate, momentum, ...)
- Optimize performance on validation set





Adaptive Computation Time for Recurrent Neural Networks

- Standard RNN: Compute one output per input
- Amount of computation controlled by network depth, the same for each input
- ACT: RNN can decide how long to cycle before producing the output
- Penalized for taking too long



Coding Session

<http://bit.ly/deeplearntut>

Coding Session with **Keras** Deep Learning Library for Python

1. Car Recognition on Images
 - Full vs. Convolutional Network
 - ReLU Activation
 - Dropout
 - Data Augmentation
2. Music vs. Speech Classification
 - Preprocessing of Audio files
 - Train/Test set split
 - single, stacked and parallel CNNs
 - RNN
3. Hard-Disk Failure Prediction with SMART data
 - RNN (*postponed to next meetup*)



Call for Meetup Topics / Talks

NEXT MEETUP: February 21, 2017 @ bwin

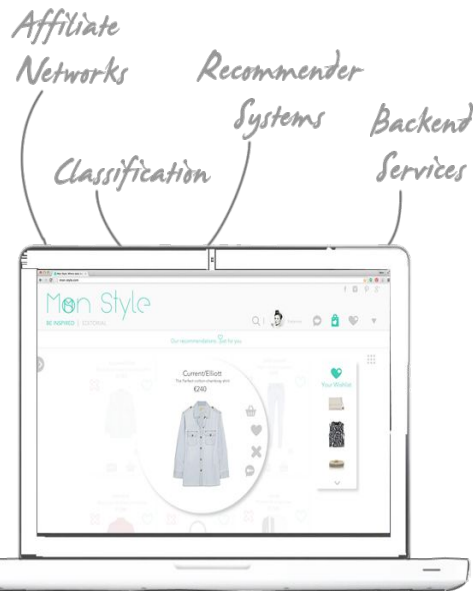

Looking for:

- Industry Applications of DL
- Real World Use Cases of DL
- Interesting novel scientific results
- Fun applications
- Niche applications
- Latest News / Hot topics

Get in touch if you want to contribute something!

tom.lidy@gmail.com

Mon Style



Job Offer in DL Startup

Mon Style is a tech startup that uses **machine-learning** to deliver highly **accurate results to online shoppers**. Their mission is to revolutionize the way we discover, interact and shop products.

Job Profile:

- Full time position based in Vienna
- Flexible working hours and home office
- Knowledge of algorithms and data structures, Python
- Knowledge of machine learning (Classification, Regression and Clustering)
- Experience using deep learning frameworks like Caffe/Torch/Theano/TensorFlow + numpy, pandas etc.

Application: office@mon-style.com



Data Science und Deep Learning Innovationslehrgang



Lehrgang (~Master class) mit 4 Modulen:

1. Data Science Basics
2. Advanced Data Science
3. Deep Learning
4. Word Embedding – Deep Learning für Text
5. Grundlagen der Teilnahme in einem Datenmarkt

durchgeführt von on TU Wien, Donau-Universität Krems und Research Studios Austria

Für Firmen:

- Beitrag zur (wirtschaftsorientierten) Ausrichtung des Lehrgangs
- **Geförderte Teilnahme am Lehrgang**
- Möglichkeit für kleine Projekte mit den Forschungspartnern
- Förderung für Teilnahme: K/M/G Unternehmen: **70/60/50 %**

Nur Firmenpartner, die Teil des Projektantrags sind, dürfen am Lehrgang teilnehmen! - Get in touch with Tom!



Deep Learning

Thanks you for coming!