Deep Learning for Object Detection in Video Surveillance

Michelangelo Fiore & Florian Matusek

KiwiSecurity
AUTOMATING VIDEO SURVEILLANCE

Number of video surveillance cameras world-wide

300m+

Incidents missed by security operators after 20 minutes

95%


VENDARE


KiwiSecurity

So where do we go from here?

Background subtraction

Background subtraction

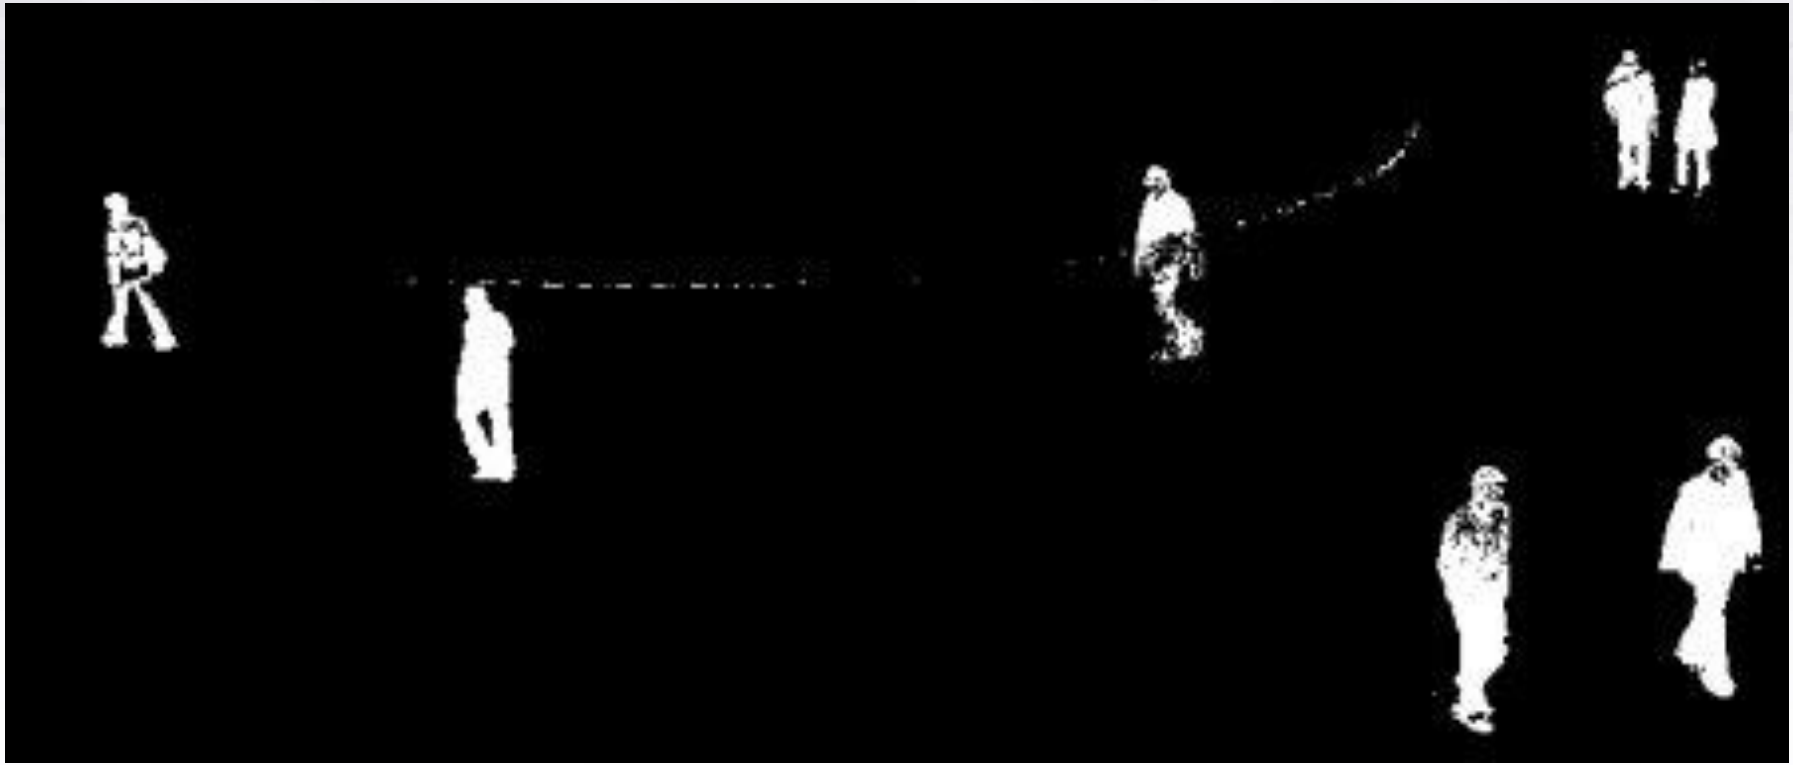

Some Pitfalls

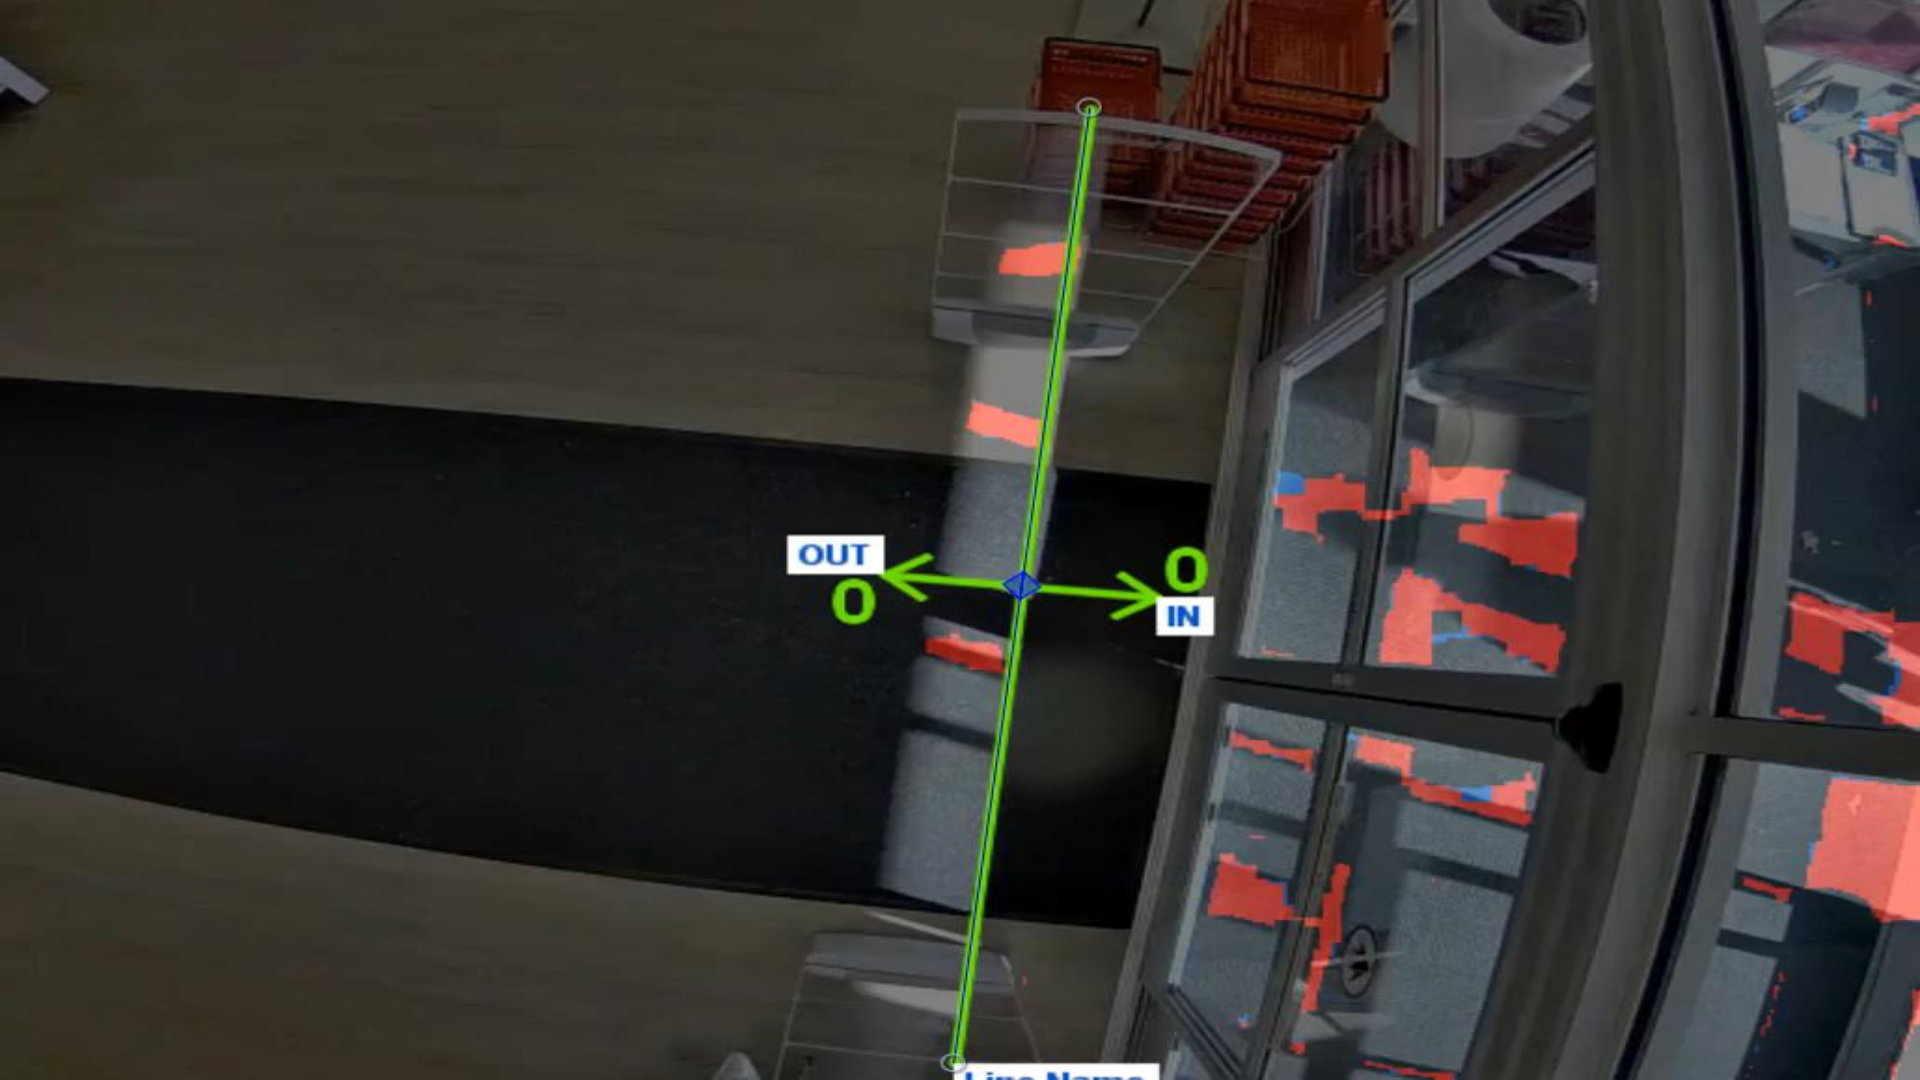
OUT
0
0
IN

KiwiSecurity

Many Parameters

Hard to Tune

Too Much

KiwiSecurity

Machine Learning Approaches

Dog: 0.99

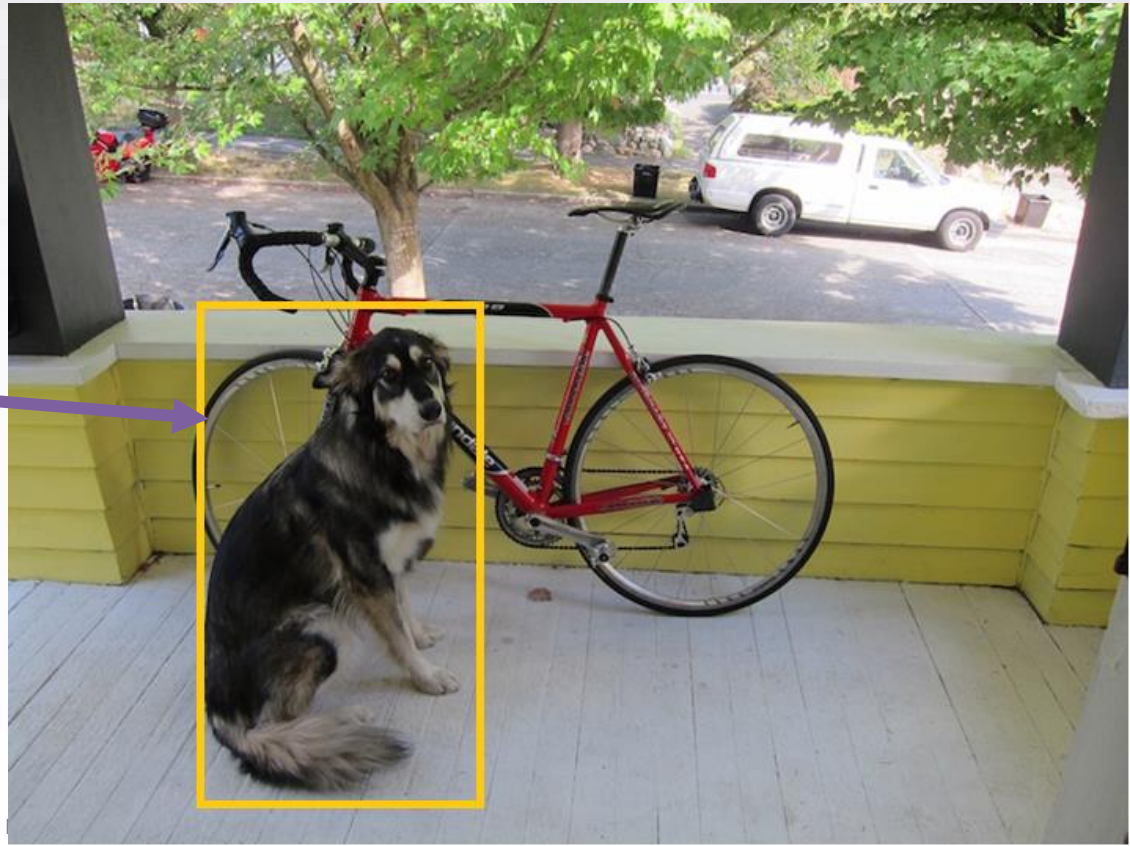

Choosing the Right Features

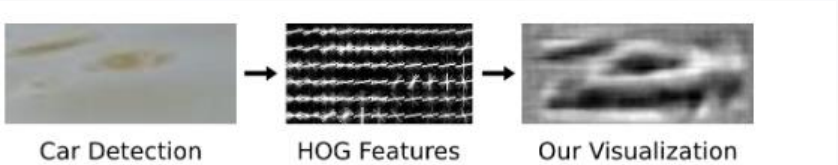

C. Vondrick, A. Khosla, T. Malisiewicz, A. Torralba. "HOGgles: Visualizing Object Detection Features" *International Conference on Computer Vision* (ICCV), Sydney, Australia, December 2013.

KiwiSecurity

Feature Extraction With Deep Learning

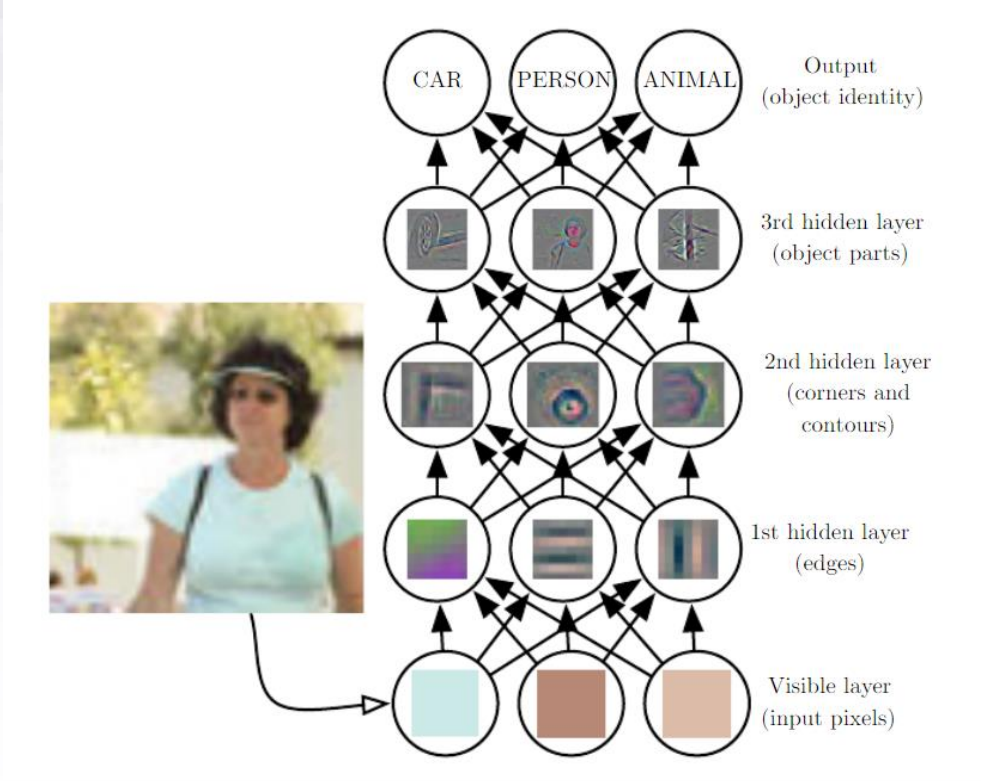

Object Detection with Deep Learning: Region Based Approach

KiwiSecurity

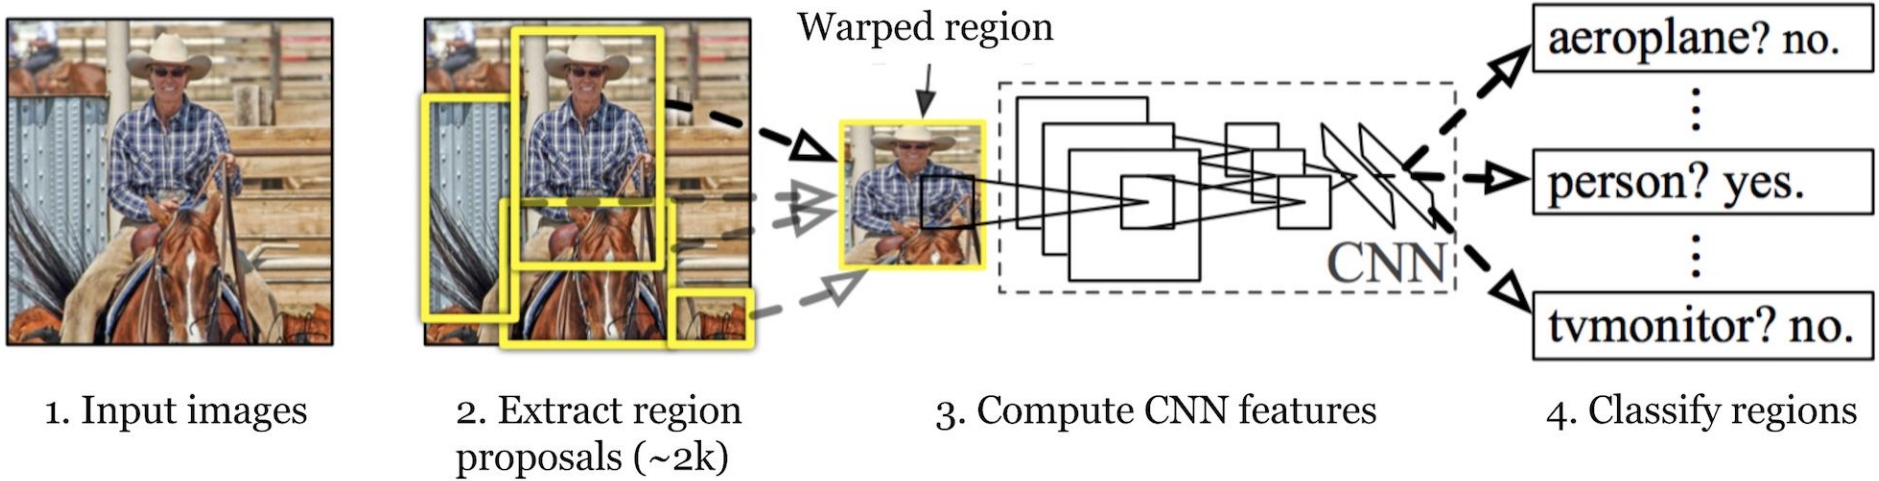

Girshick, Ross, et al. "Rich feature hierarchies for accurate object detection and semantic segmentation." *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2014.

Object Detection with Deep Learning: Single Shot Approach

KiwiSecurity

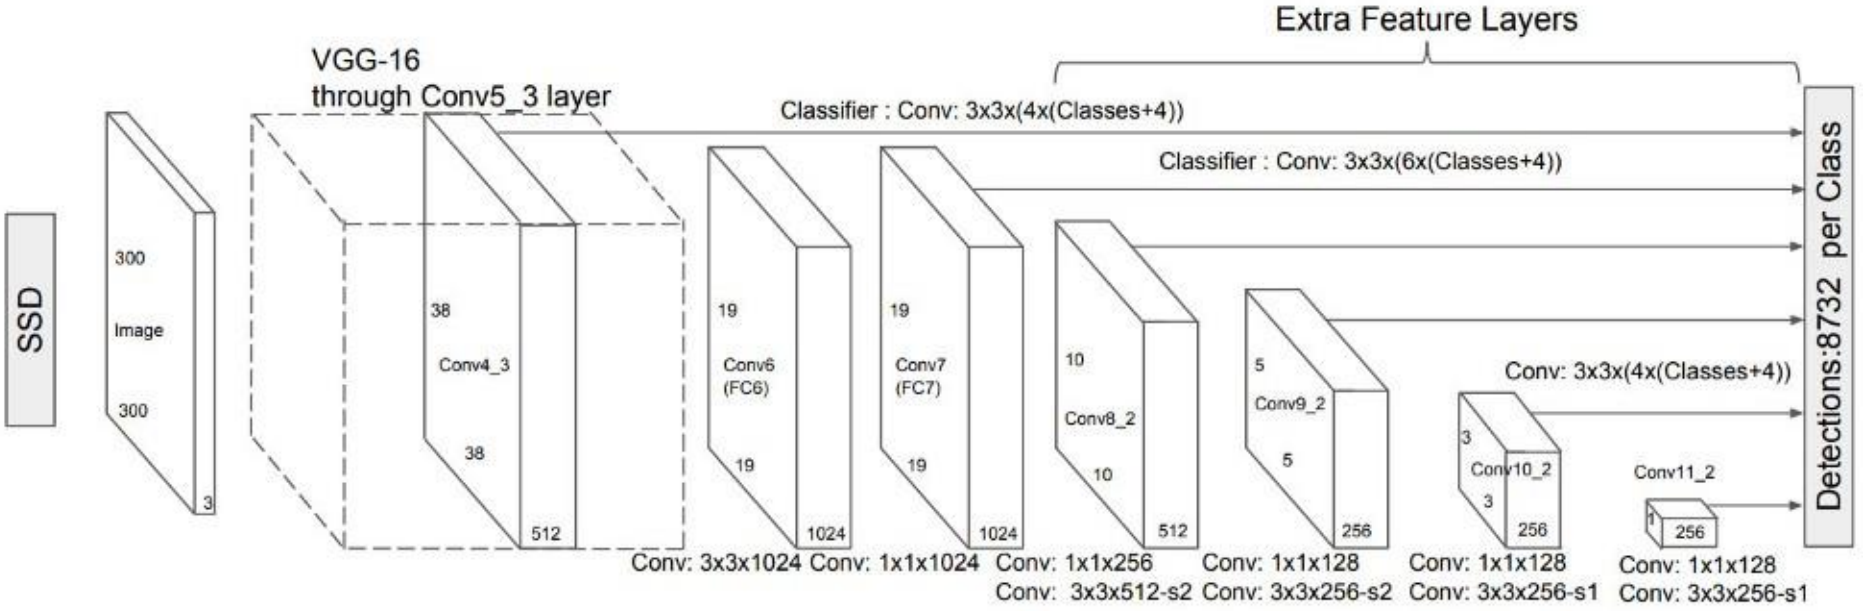

Liu, Wei, et al. "Ssd: Single shot multibox detector." *European conference on computer vision*. Springer, Cham, 2016

Detection at Different Scales

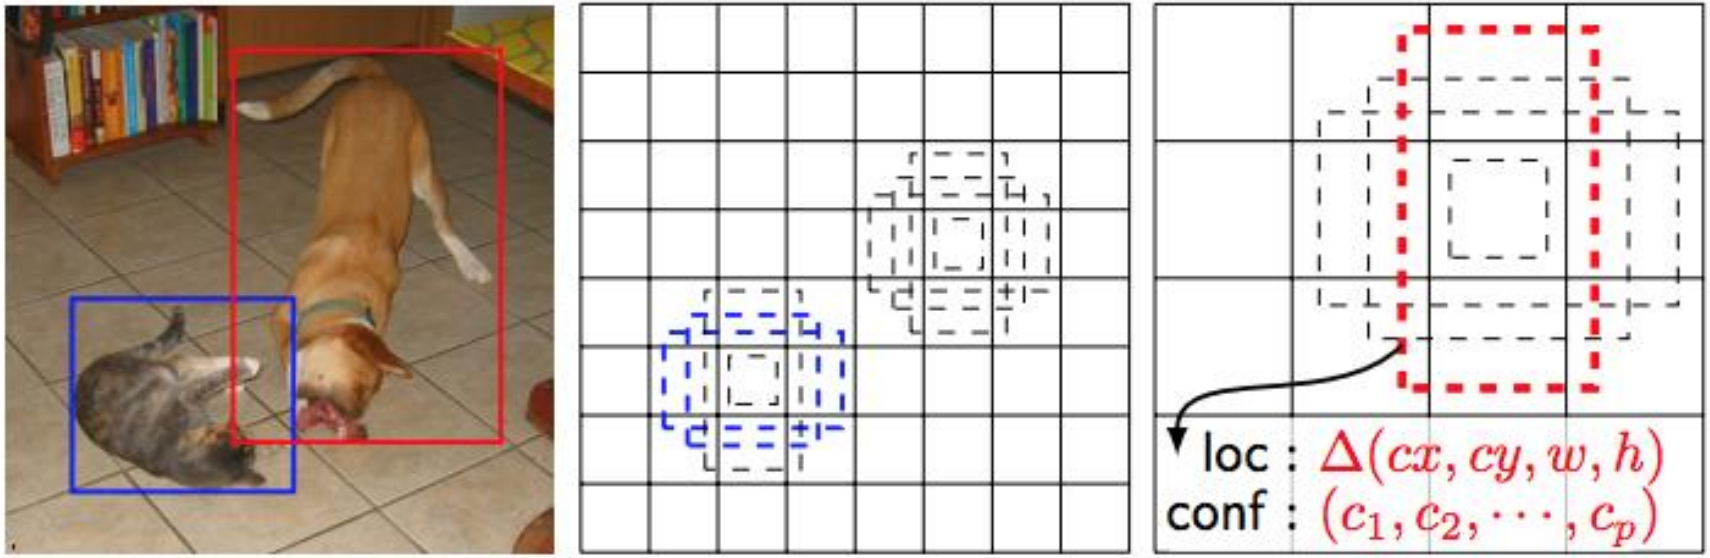

(a) Image with GT boxes (b) 8×8 feature map (c) 4×4 feature map

Choosing a Detector

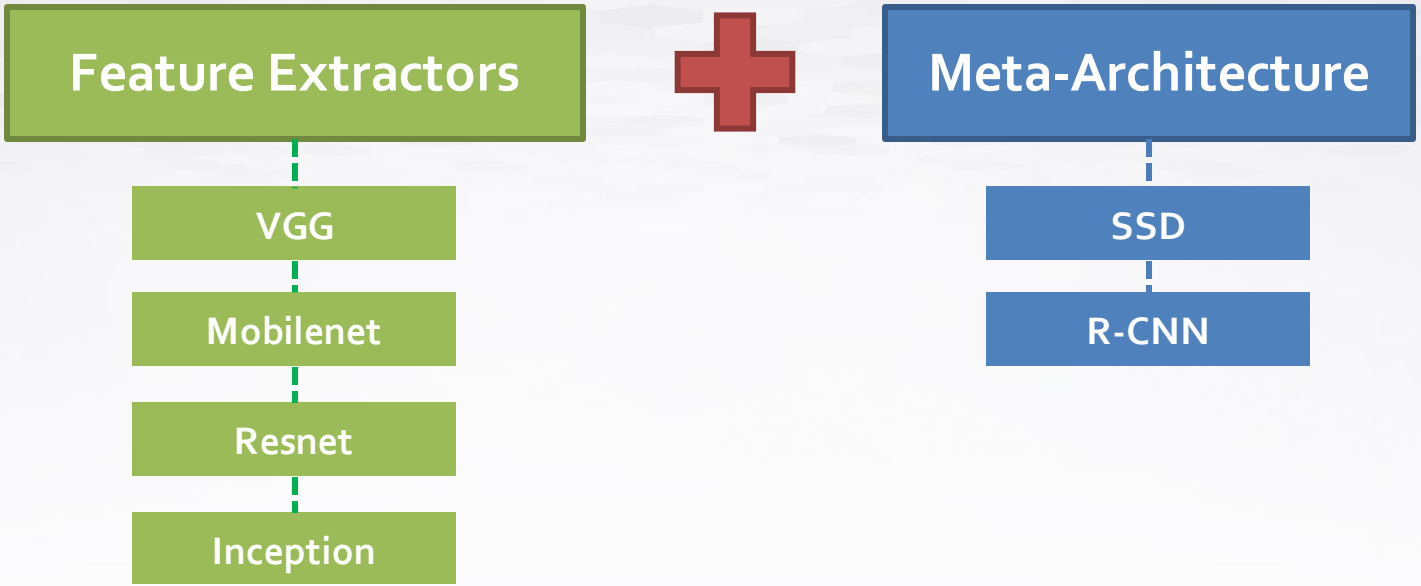

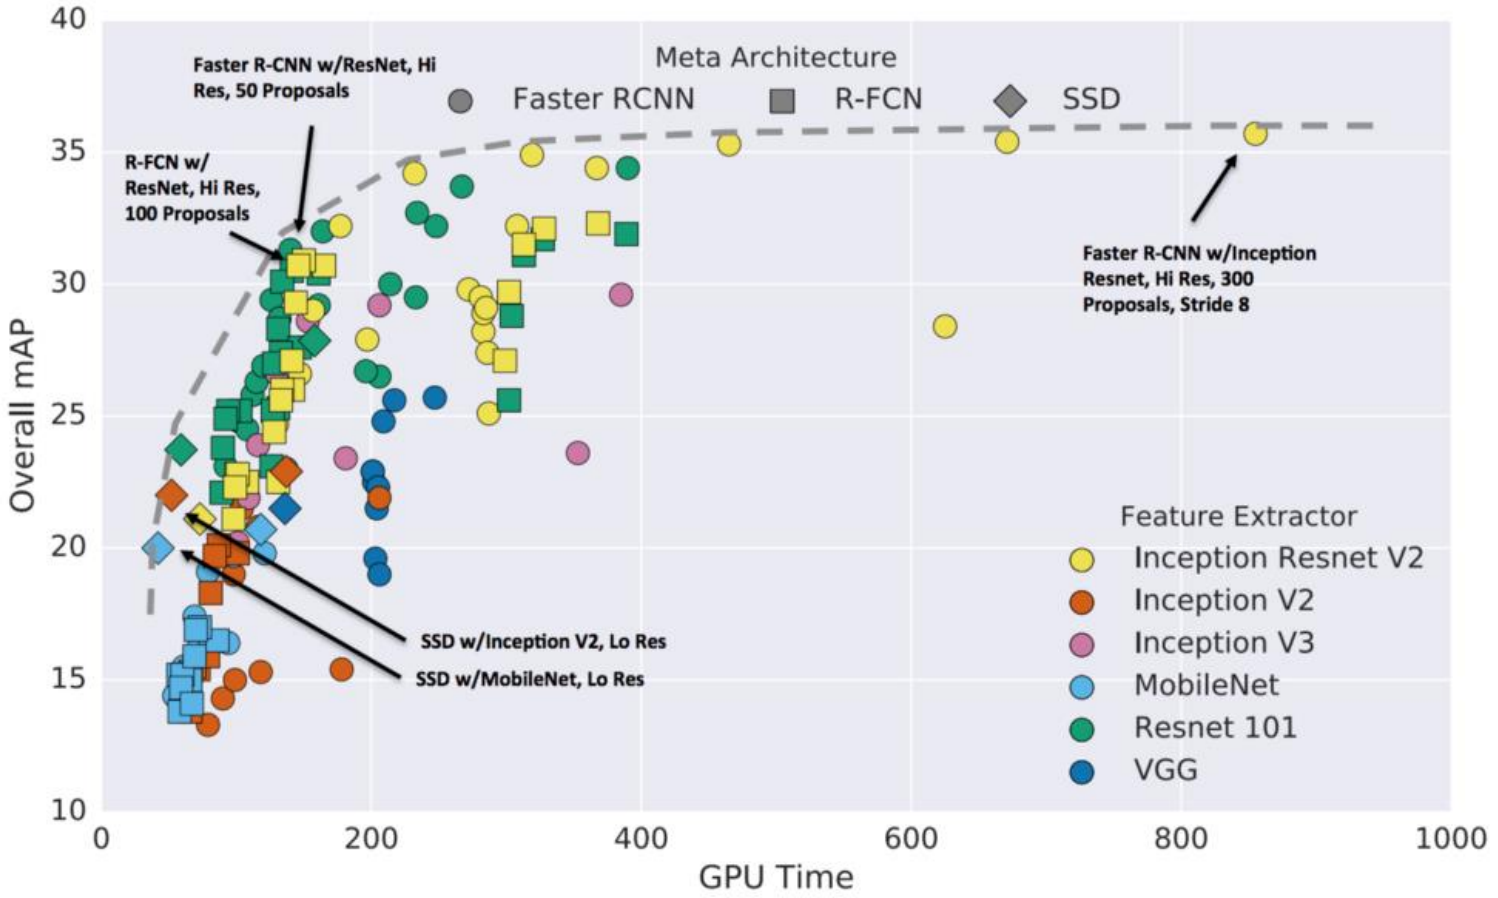

Huang, Jonathan, et al. "Speed/accuracy trade-offs for modern convolutional object detectors." *IEEE CVPR*. 2017.

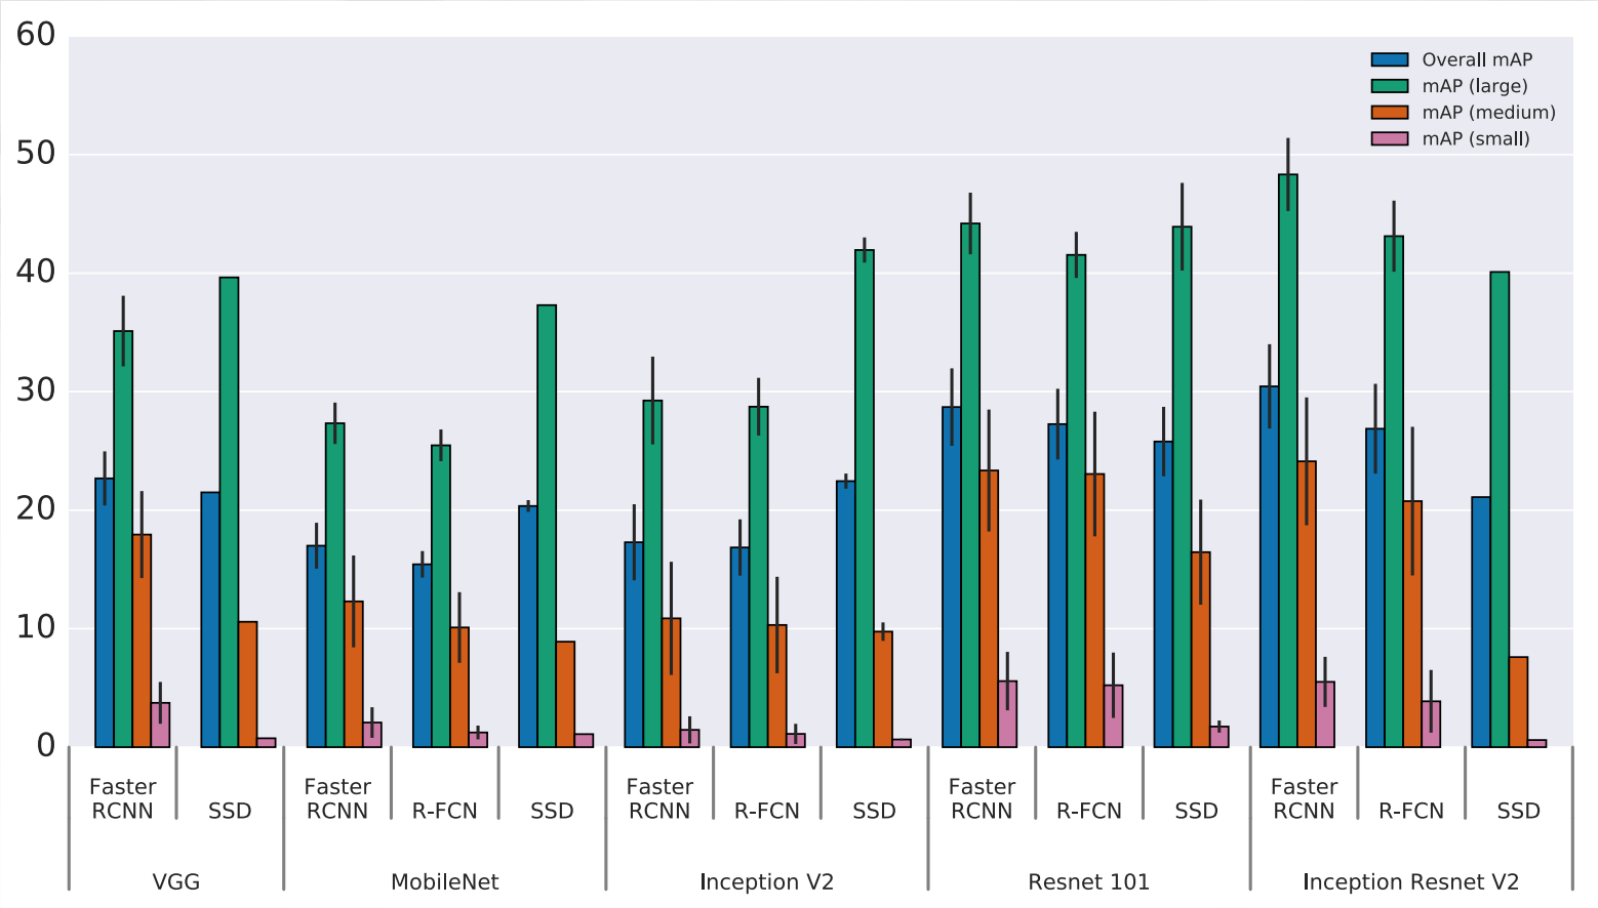
60
50
40
30
20
10
0
Overall mAP
mAP (large)
mAP (medium)
mAP (small)
Faster RCNN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
Faster RCNN
R-FCN
SSD
VGG
MobileNet
Inception V2
Resnet 101
Inception Resnet V2

Understanding Your Requirements

Our Frameworks: Apache MXNet

OS and Language support	
Documentation	Good API and GLUON
Models Support	Many Object Detection Models
Performances	Fast Performances on Training and Inference

https://mxnet.apache.org

Our Frameworks: Intel OpenVINO

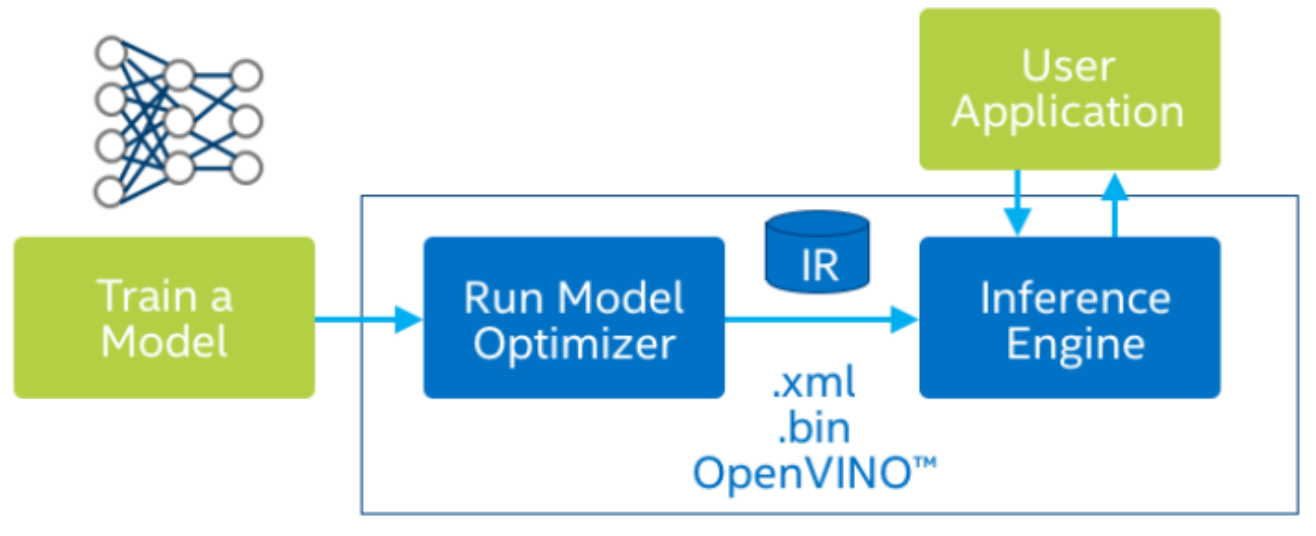

https://software.intel.com/en-us/openvino-toolkit

Our Detector

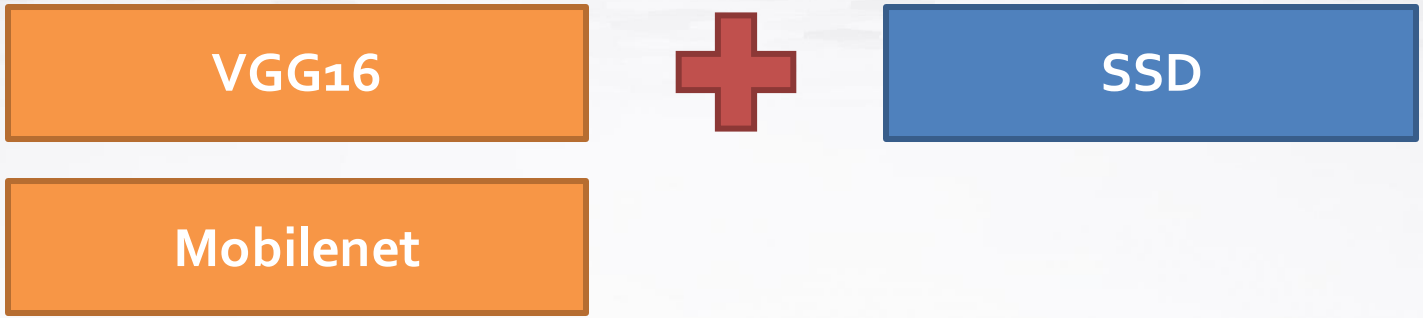

Building a Dataset

KiwiSecurity

People Detection Dataset

KiwiSecurity

Some Results

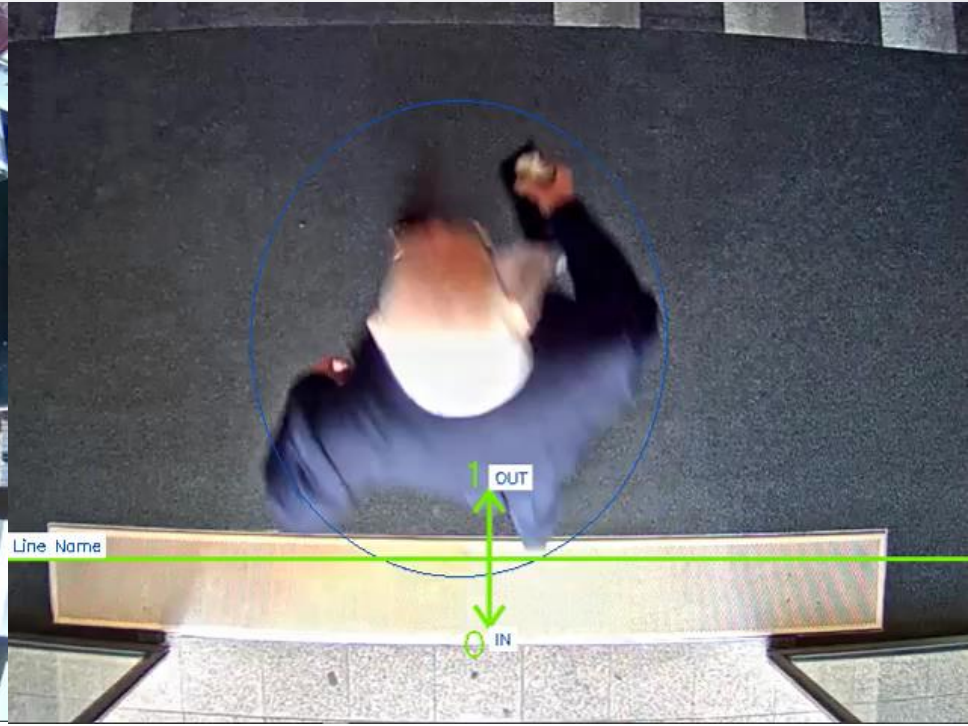

MainWindow

File Help

Expected 57

Frame 23; 00:00:01.920
Playing
AfterProcess

ITA 1:16 PM 2018-10-02

Improving the Performances
Reducing the Model Complexity

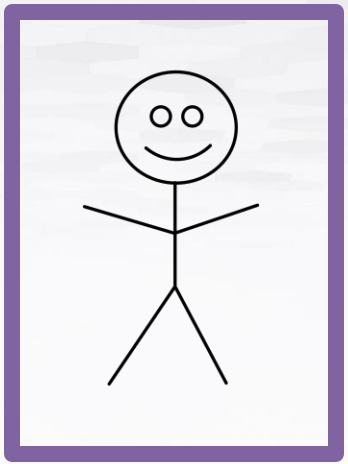

Model Pruning

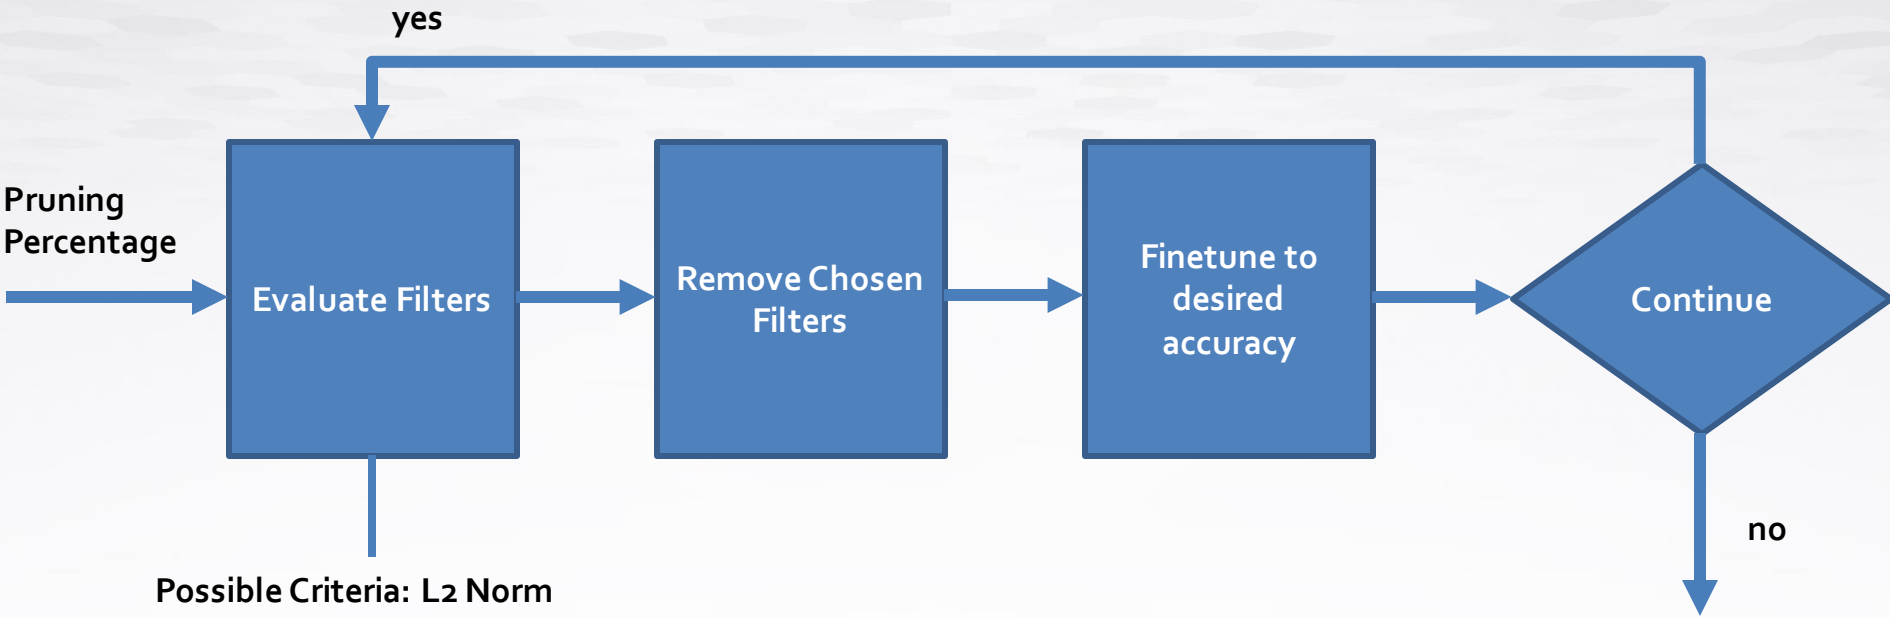

Molchanov, Pavlo, et al. "Pruning convolutional neural networks for resource efficient inference." (2016).

Pruning Results: VGG16

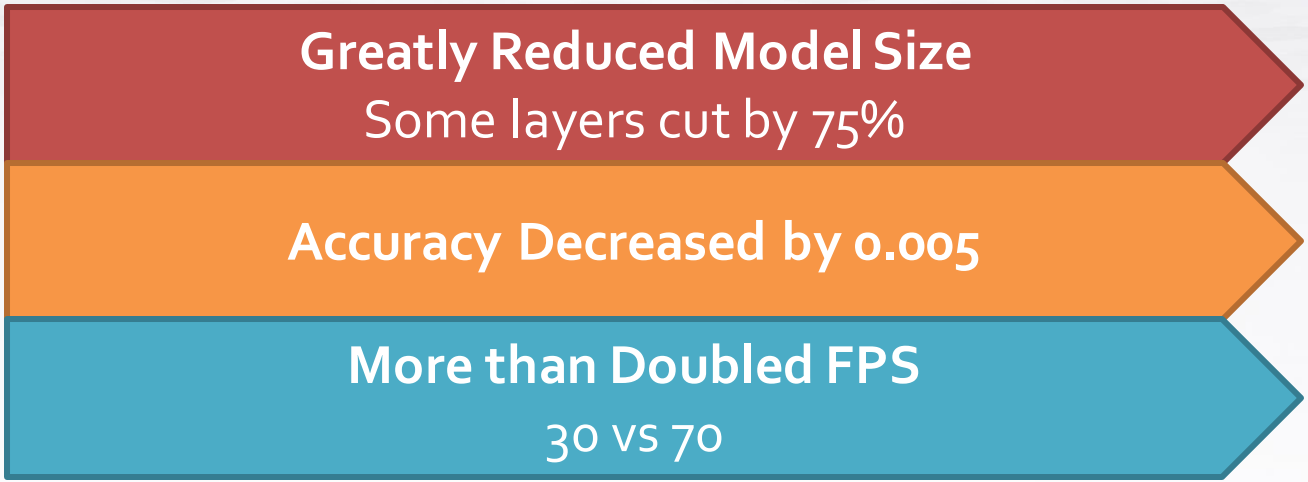

All experiments where done on an Intel Core i7-7800X CPU with a Nvidia Quadro P2000 GPU

Improving the Performances
Do We Need to Detect at Every Frame?

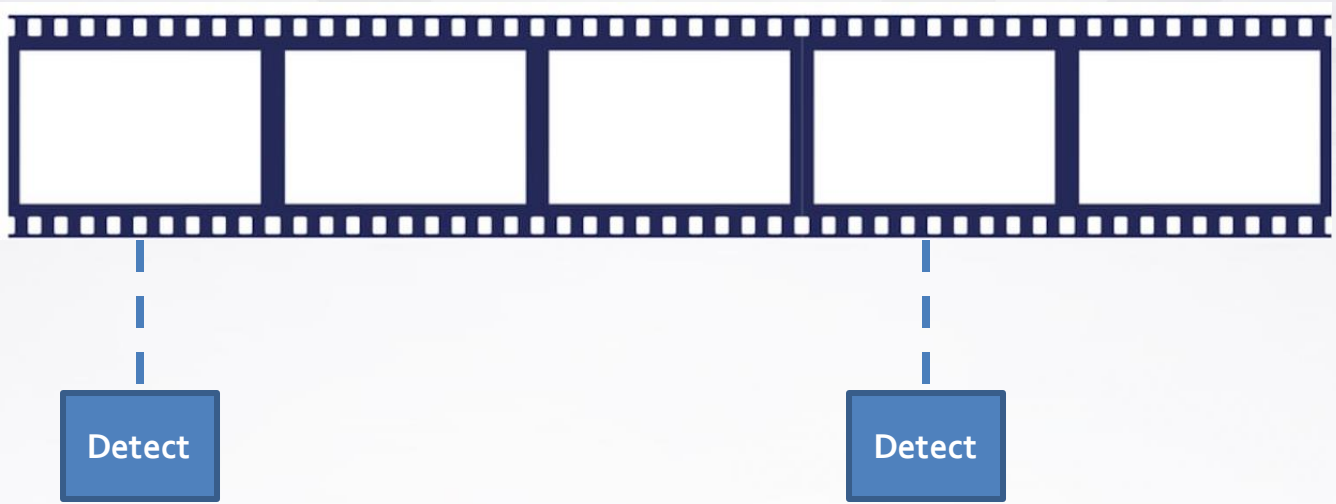

Improving the Performances
Reducing the Detection Area

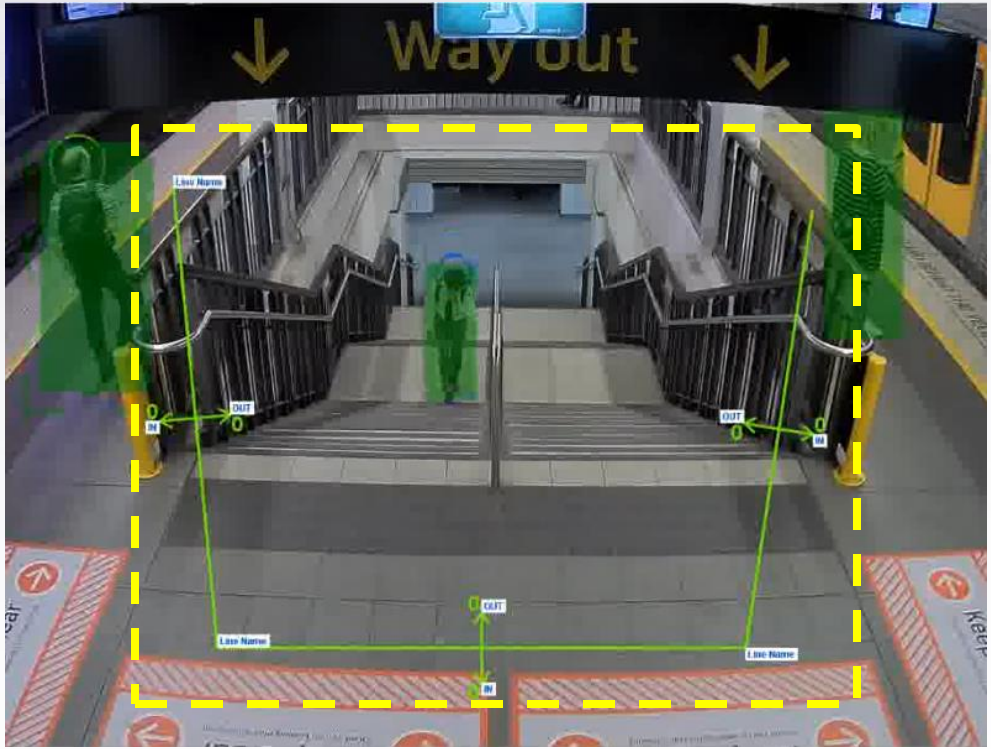

1	Reduce Distorsion
2	Increase Object Size

Significant Performance Increase (6%)

If you don't need to beat ImageNet, **don't try to**

Use **every trick** you can

Future Plans

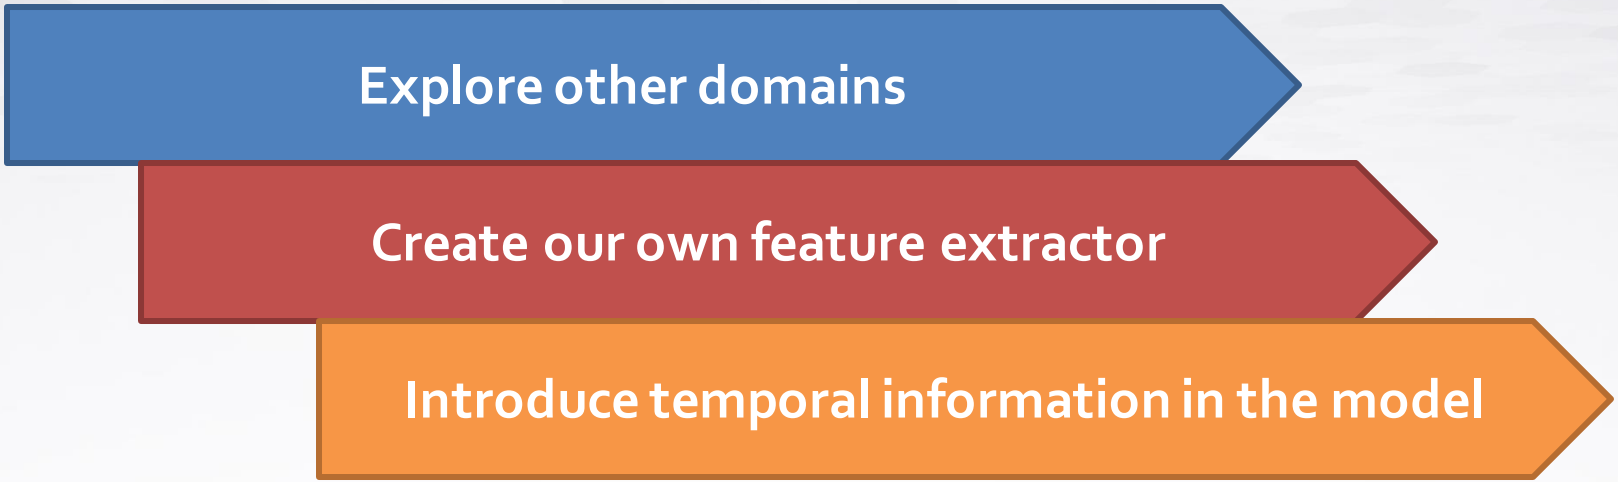

Interested?

Michelangelo Fiore
m.fiore@kiwisecurity.com

Florian Matusek
f.matusek@kiwisecurity.com

www.kiwisecurity.com

PS: We're hiring!